

Class: Final Year (Computer Science and Engineering)

Year: 2024-25

Semester: 1

Course: High Performance Computing Lab

Practical No. 5

Exam Seat No: 22510021

Title of practical: Implementation of OpenMP programs.

Implement following Programs using OpenMP with C:

Problem Statement 1: Implementation of Matrix-Matrix Multiplication.

Screenshots:

```
PS C:\Lab> cd hpc\5
PS C:\Lab\hpc\5> gcc -fopenmp 1Matrix.c -o matrix
PS C:\Lab\hpc\5> .\matrix
Matrix multiplication completed.
Time taken: 0.002000 seconds
C[0][0] = 0, C[0][1] = 0, C[1][0] = 0
PS C:\Lab\hpc\5>
```

Information:

Program Steps:

1. Declare three $N \times N$ matrices: A, B, and C.
2. Initialize $A[i][j] = i$, $B[i][j] = j$, and $C[i][j] = 0$.
3. Set number of threads using `omp_set_num_threads(4)`.
4. Use OpenMP parallel for with `collapse(2)` to compute each element of $C[i][j]$ in parallel.
5. Measure start and end time using `omp_get_wtime()`.
6. Print the total time taken and a few sample elements of C to verify correctness.

Analysis:

- Threads compute complete elements independently \rightarrow no race conditions.
- Example (2 threads, 4×4 matrix):
 - Thread 0: $C[0][0]$, $C[0][2]$, $C[1][0]$, $C[1][2]$

- Thread 1: C[0][1], C[0][3], C[1][1], C[1][3]
- Parallel execution reduces time significantly.
- Small matrices may show less improvement due to thread overhead.
- Increasing threads improves speedup up to the number of CPU cores.

Problem Statement 2: Implementation of Matrix-scalar Multiplication.

Screenshots:

```
PS C:\Lab\hpc\5> gcc -fopenmp 2scaler.c -o scaler
PS C:\Lab\hpc\5> .\scaler
PS C:\Lab\hpc\5> gcc -fopenmp 2scaler.c -o scaler
PS C:\Lab\hpc\5> .\scaler
Scalar multiplication completed.
Time taken: 0.012000 seconds
C[0][0] = 0, C[0][1] = 0, C[1][0] = 5
PS C:\Lab\hpc\5> █
```

Information:

- Scalar Multiplication Formula: $C[i][j] = \text{scalar} \times A[i][j]$
- OpenMP Parallelization:
 - `#pragma omp parallel for collapse(2)` → Distributes the row-column iterations among multiple threads.
 - `omp_set_num_threads()` → Sets the number of threads globally.
- `omp_get_wtime()` → Measures the wall-clock time for execution.

Program Steps:

1. Declare matrices A and C of size $N \times N$.
2. Initialize matrix A with values.
3. Set the number of threads using `omp_set_num_threads()`.
4. Multiply each element by the scalar using OpenMP parallel loops.
5. Measure and print the time taken and a few elements of the result matrix.

Analysis:

Time Complexity:

- Sequential: $O(N^2)O(N^2)O(N^2)$

- Parallel: The **i** and **j** loops are split among threads, so effective time decreases roughly by the number of threads.

Thread Work Distribution:

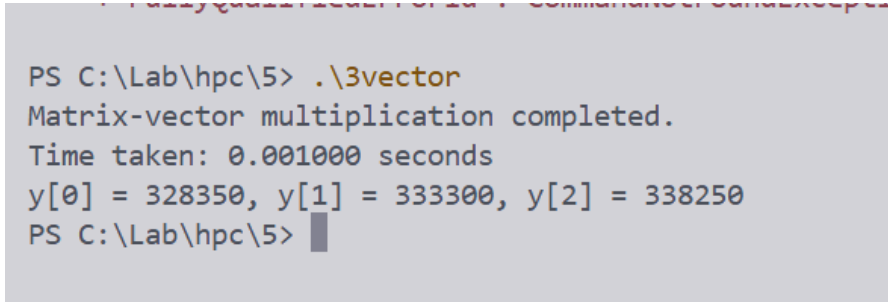
- Using collapse 2 threads are assigned blocks of i and j pairs.
- Each thread multiplies its assigned elements independently → no race conditions.

Observations:

- Parallel execution significantly reduces computation time, especially for large matrices.
- Small matrices may show minor improvement due to thread overhead.
- Increasing threads improves speedup up to the number of CPU cores.

Problem Statement 3: Implementation of Matrix-Vector Multiplication.

Screenshots:



```
PS C:\Lab\hpc\5> .\3vector
Matrix-vector multiplication completed.
Time taken: 0.001000 seconds
y[0] = 328350, y[1] = 333300, y[2] = 338250
PS C:\Lab\hpc\5>
```

Information:

Program Steps:

1. Declare a matrix $A[N][N]$, input vector $x[N]$, and result vector $y[N]$.
2. Initialize matrix and vector with values.
3. Set the number of threads using `omp_set_num_threads()`.
4. Parallelize the computation of each $y[i]$ using `#pragma omp parallel for`.
5. Compute each element as the dot product of a row of A and vector x .
6. Measure execution time using `omp_get_wtime()` and print sample elements of y .

Analysis:

- $O(N^2)$

- Parallel: Each row computation is independent, so rows are distributed among threads for faster execution.
- Each thread computes a subset of rows completely.
- No race conditions occur because each thread writes to its own element of y.

Problem Statement 4: Implementation of Prefix sum.

Screenshots:

```
PS C:\Lab\hpc\5> .\3vector
Matrix-vector multiplication completed.
Time taken: 0.001000 seconds
y[0] = 328350, y[1] = 333300, y[2] = 338250
PS C:\Lab\hpc\5> gcc -fopenmp 4prefixsum.c -o prefixsum
PS C:\Lab\hpc\5> .\prefixsum
Prefix sum completed.
Time taken: 0.012000 seconds
Array: 1 3 6 10 15 21 28 36 45 55
PS C:\Lab\hpc\5>
```

Ln 40,

Information:

Program Steps:

1. Declare input array A[N] and result array B[N].
2. Initialize A with values.
3. Set number of threads using `omp_set_num_threads()`.
4. Compute prefix sum in parallel (simple version).
5. Measure execution time using `omp_get_wtime()` and print the result array.

Analysis:

1. Time Complexity: Sequential $O(N)$
2. Parallel: Basic version has overhead; optimal version achieves $O(N/p + \log p)$, where p = number of threads.
3. Work Distribution :
 - a. Each thread computes partial sums of assigned chunk.
 - b. Then, offsets are added to adjust the sums across threads.

Github Link: <https://github.com/22510021-Shrikrishna/HPC.git>